



The beauty of trees

Statistical Methods for Data Science

Federico Pigozzi

Letícia Negrão Pinto

Eduardo Gonnelli

DSSC - 2019

tree R package

Classification trees
Regression trees

Package ‘tree’

April 26, 2019

Title Classification and Regression Trees

Version 1.0-40

Date 2019-03-01

Depends R (>= 3.6.0), grDevices, graphics, stats

Suggests MASS

Description Classification and regression trees.

License GPL-2 | GPL-3

NeedsCompilation yes

Author Brian Ripley [aut, cre]

Maintainer Brian Ripley <ripley@stats.ox.ac.uk>

Repository CRAN

Date/Publication 2019-04-26 13:21:21 UTC

R topics documented:

cv.tree	2
deviance.tree	3
misclass.tree	3
na.tree.replace	4
partition.tree	5
plot.tree	6
plot.tree.sequence	7
predict.tree	8
prune.tree	9
snip.tree	11
text.tree	12
tile.tree	13
tree	14
tree.control	16
tree.screens	17

Index

19

Regression trees - Boston data set

```
# sample rows for the training set, choose a 25% hold-out validation set  
train <- sample(1: nrow(data), nrow(data) * 0.75)  
boston.test <- data[-train , !(names(data) %in% c("medv"))] # omit "medv" as it is the target  
y.true <- data[-train, "medv"]
```

Building a Tree

- Recursive binary splitting
- Impurity measures: deviance (default), Gini

```
tree.boston <- tree(medv~., data, subset=train) # all variables, subset train
```

```
pred.tree <- predict(tree.boston, newdata=boston.test, type="vector")  
RMSE(pred.tree, y.true)
```

```
## [1] 4.262707
```

Visualize Tree Information - 1

```
summary(tree.boston)
```

```
##
```

```
## Regression tree:
```

```
## tree(formula = medv ~ ., data = data, subset = train)
```

```
## Variables actually used in tree construction:
```

```
## [1] "rm"      "lstat"    "age"      "crim"      "ptratio"
```

```
## Number of terminal nodes: 9
```

```
## Residual mean deviance: 14.55 = 5383 / 370
```

```
## Distribution of residuals:
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
## -14.8000 -2.0910 -0.1253  0.0000  2.0730 17.1700
```

Visualizing Tree Information - 2

```
tree.boston
```

```
## node), split = , deviance, yval
##      * denotes terminal node
##
##  1) root 379 30990.0 22.63
##    2) rm < 6.797 307 10630.0 19.66
##      4) lstat < 14.405 179 3665.0 22.87
##        8) lstat < 9.715 92 2322.0 24.89
##          16) age < 89.45 87 965.2 24.25 *
##            17) age > 89.45 5 702.4 36.02 *
##
##          9) lstat > 9.715 87 565.7 20.73 *
##        5) lstat > 14.405 128 2551.0 15.18
##          10) crim < 6.91188 77 922.7 17.48 *
##            11) crim > 6.91188 51 605.5 11.70 *
##    3) rm > 6.797 72 6076.0 35.31
##      6) rm < 7.47 51 2299.0 31.27
##        12) crim < 7.39342 46 996.9 32.83 *
##          13) crim > 7.39342 5 166.9 16.96 *
##    7) rm > 7.47 21 923.2 45.12
##      14) ptratio < 17.6 16 139.0 47.76 *
##        15) ptratio > 17.6 5 318.5 36.70 *
```

Visualizing Tree Information - 2

```
tree.boston
```

```
## node), split = , deviance = ,
##      * dev # observations
##
## 1) root 379 30990.0 22.63
##    2) rm < 6.797 307 10630.0 19.66
##      4) lstat < 14.405 179 3665.0 22.87
##        8) lstat < 9.715 92 2322.0 24.89
##          16) age < 89.45 87 965.2 24.25 *
##            17) age > 89.45 5 702.4 36.02 *
##
##          9) lstat > 9.715 87 565.7 20.73 *
##        5) lstat > 14.405 128 2551.0 15.18
##      10) crim < 6.91188 77 922.7 17.48 *
##      11) crim > 6.91188 51 605.5 11.70 *
##    3) rm > 6.797 72 6076.0 35.31
##      6) rm < 7.47 51 2299.0 31.27
##        12) crim < 7.39342 46 996.9 32.83 *
##        13) crim > 7.39342 5 166.9 16.96 *
##      7) rm > 7.47 21 923.2 45.12
##        14) ptratio < 17.6 16 139.0 47.76 *
##        15) ptratio > 17.6 5 318.5 36.70 *
```

Visualizing Tree Information - 2

```
tree.boston
```

```
## node), split = deviance, yval
##      * deviance node
##
## 1) root 379 30990.0 22 63
##    2) rm < 6.797 307 10630.0 19.66
##      4) lstat < 14.405 179 3665.0 22.87
##        8) lstat < 9.715 92 2322.0 24.89
##          16) age < 89.45 87 965.2 24.25 *
##            17) age > 89.45 5 702.4 36.02 *
##
##          9) lstat > 9.715 87 565.7 20.73 *
##        5) lstat > 14.405 128 2551.0 15.18
##          10) crim < 6.91188 77 922.7 17.48 *
##            11) crim > 6.91188 51 605.5 11.70 *
##    3) rm > 6.797 72 6076.0 35.31
##      6) rm < 7.47 51 2299.0 31.27
##        12) crim < 7.39342 46 996.9 32.83 *
##          13) crim > 7.39342 5 166.9 16.96 *
##    7) rm > 7.47 21 923.2 45.12
##      14) ptratio < 17.6 16 139.0 47.76 *
##        15) ptratio > 17.6 5 318.5 36.70 *
```


Visualize Tree Information - 2

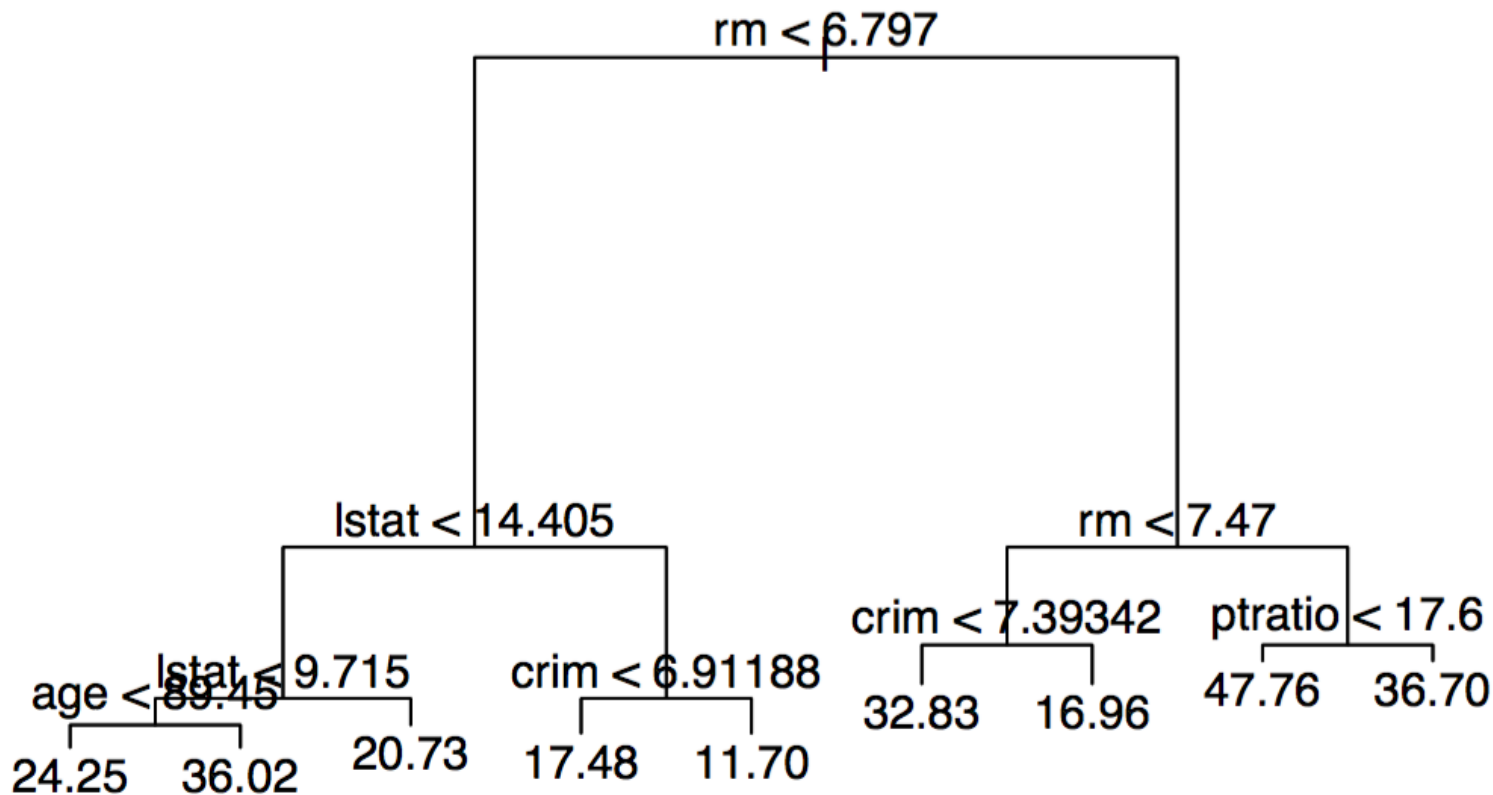
```
tree.boston
```

```
## node), split, n, deviance, yval
##      * deviance information
##
## 1) root 379 30990.0 22.63
##    2) rm < 6.797 307 10630.0 19.66
##      4) lstat < 14.405 179 3665.0 22.87
##        8) lstat < 9.715 92 2322.0 24.89
##          16) age < 89.45 87 965.2 24.25 *
##            17) age > 89.45 5 702.4 36.02 *
##
##          9) lstat > 9.715 87 565.7 20.73 *
##        5) lstat > 14.405 128 2551.0 15.18
##          10) crim < 6.91188 77 922.7 17.48 *
##            11) crim > 6.91188 51 605.5 11.70 *
##    3) rm > 6.797 72 6076.0 35.31
##      6) rm < 7.47 51 2299.0 31.27
##        12) crim < 7.39342 46 996.9 32.83 *
##          13) crim > 7.39342 5 166.9 16.96 *
##    7) rm > 7.47 21 923.2 45.12
##      14) ptratio < 17.6 16 139.0 47.76 *
##        15) ptratio > 17.6 5 318.5 36.70 *
```

mean value (class probs. for classification)

Visualize Tree Information - 3

```
# let's produce an annotated tree
plot(tree.boston)
text(tree.boston, pretty=0)
```



Building a Tree - rpart

- Open-source implementation of CART (Breiman et al. 1983)
- Same basic signature as tree

```
boston.rpart <- rpart(medv~., data, subset=train, cp=1e-05, method="anova")  
predict.rpart <- predict(boston.rpart, newdata=boston.test)  
RMSE(predict.rpart, y.true)
```

```
## [1] 4.118872
```

```
pred.tree <- predict(tree.boston, newdata=boston.test, type="vector")  
RMSE(pred.tree, y.true)
```

```
## [1] 4.262707
```

Time Comparison

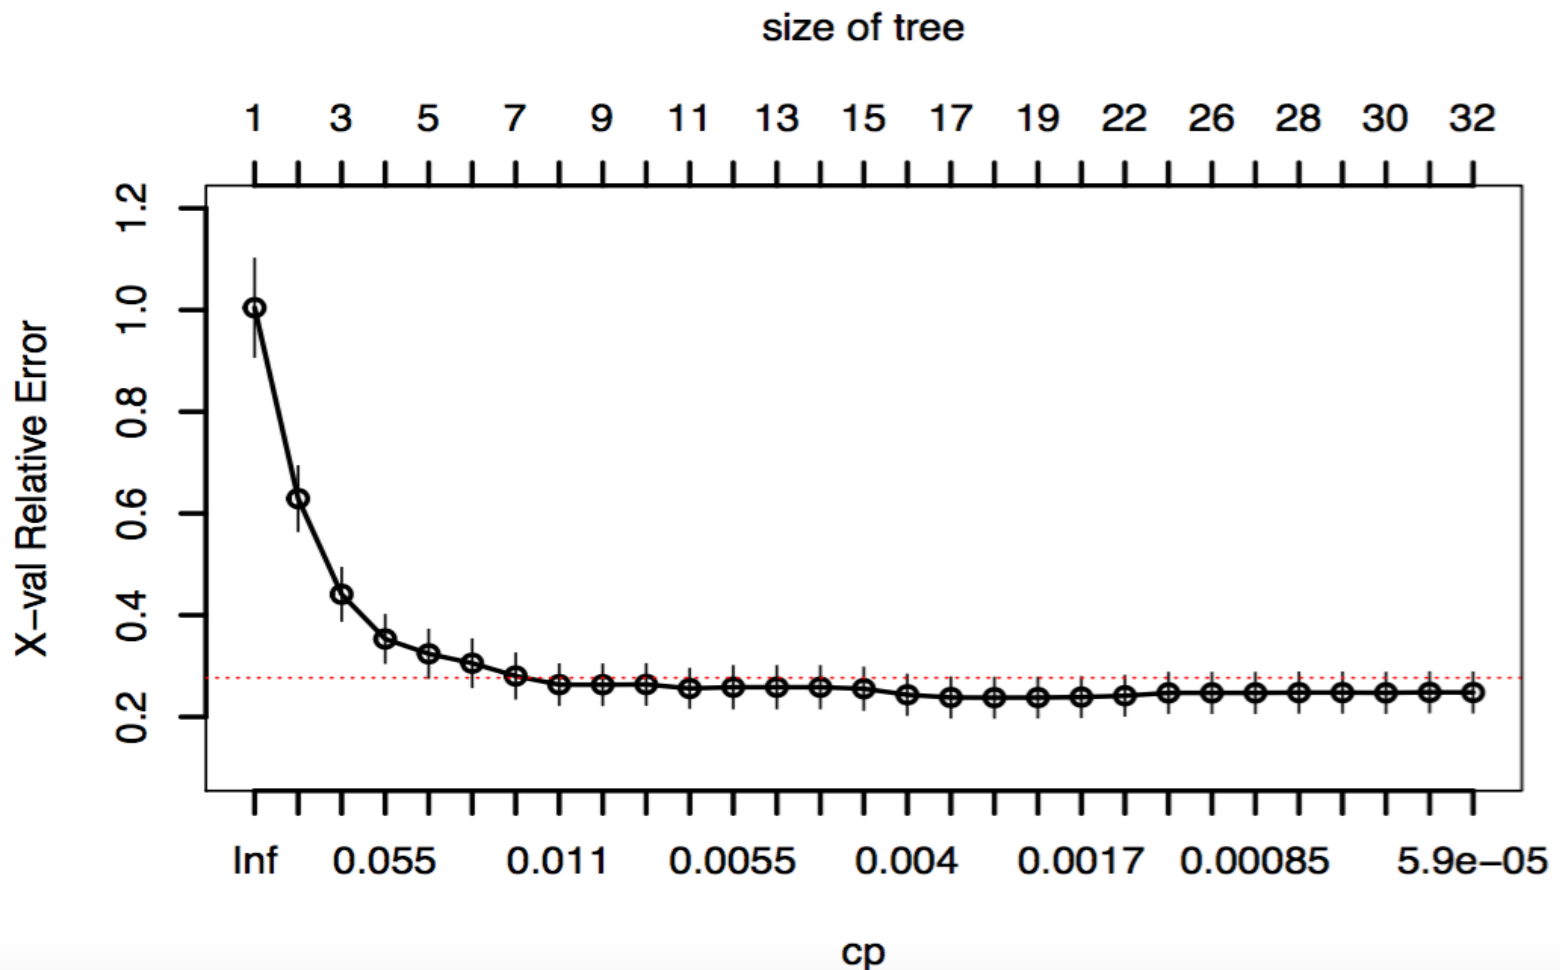
```
nsim <- 100
benchmark("grow.rpart" = {rpart(medv~., data, subset=train, cp=1e-05, method="anova")},
          "grow.tree" = {tree(medv~., data, subset=train)}),
  replications=nsim,
  columns=c("test", "replications", "elapsed", "relative", "user.self", "sys.self"))
```

##	test	replications	elapsed	relative	user.self	sys.self
## 1	grow.rpart	100	1.441	3.743	1.418	0.014
## 2	grow.tree	100	0.385	1.000	0.377	0.005

Reason: rpart performs cross-validation

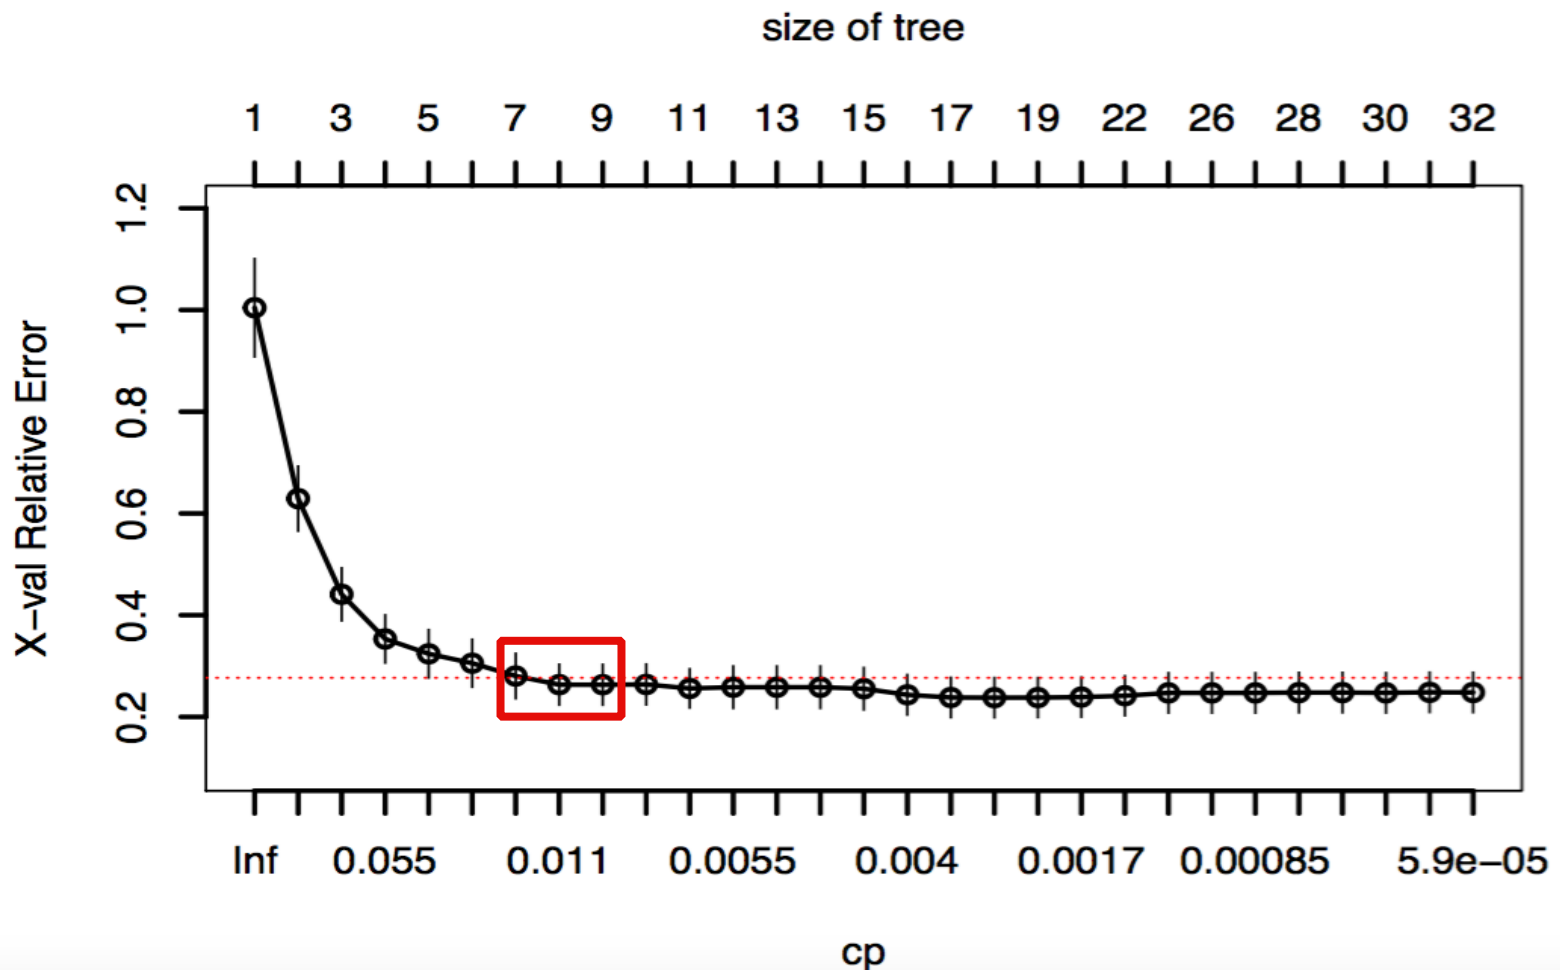
rpart – visualizing cv results

```
plotcp(boston.rpart, minline=TRUE, col="red", lwd=2)
```



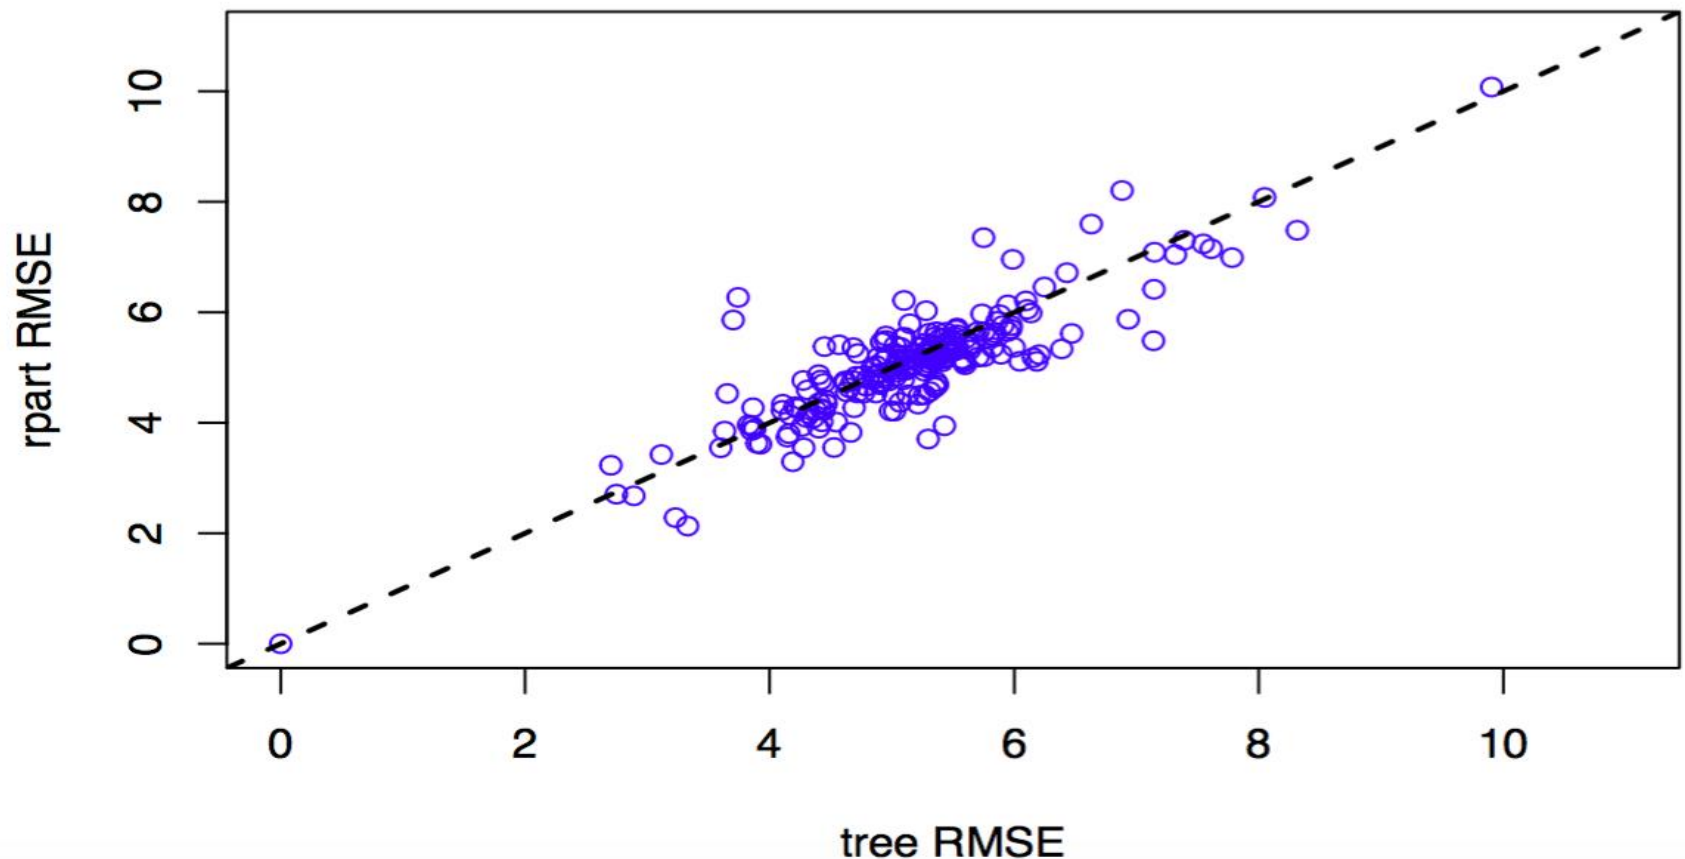
rpart – visualizing cv results

```
plotcp(boston.rpart, minline=TRUE, col="red", lwd=2)
```



Performance Comparison

RMSE for rpart vs. tree

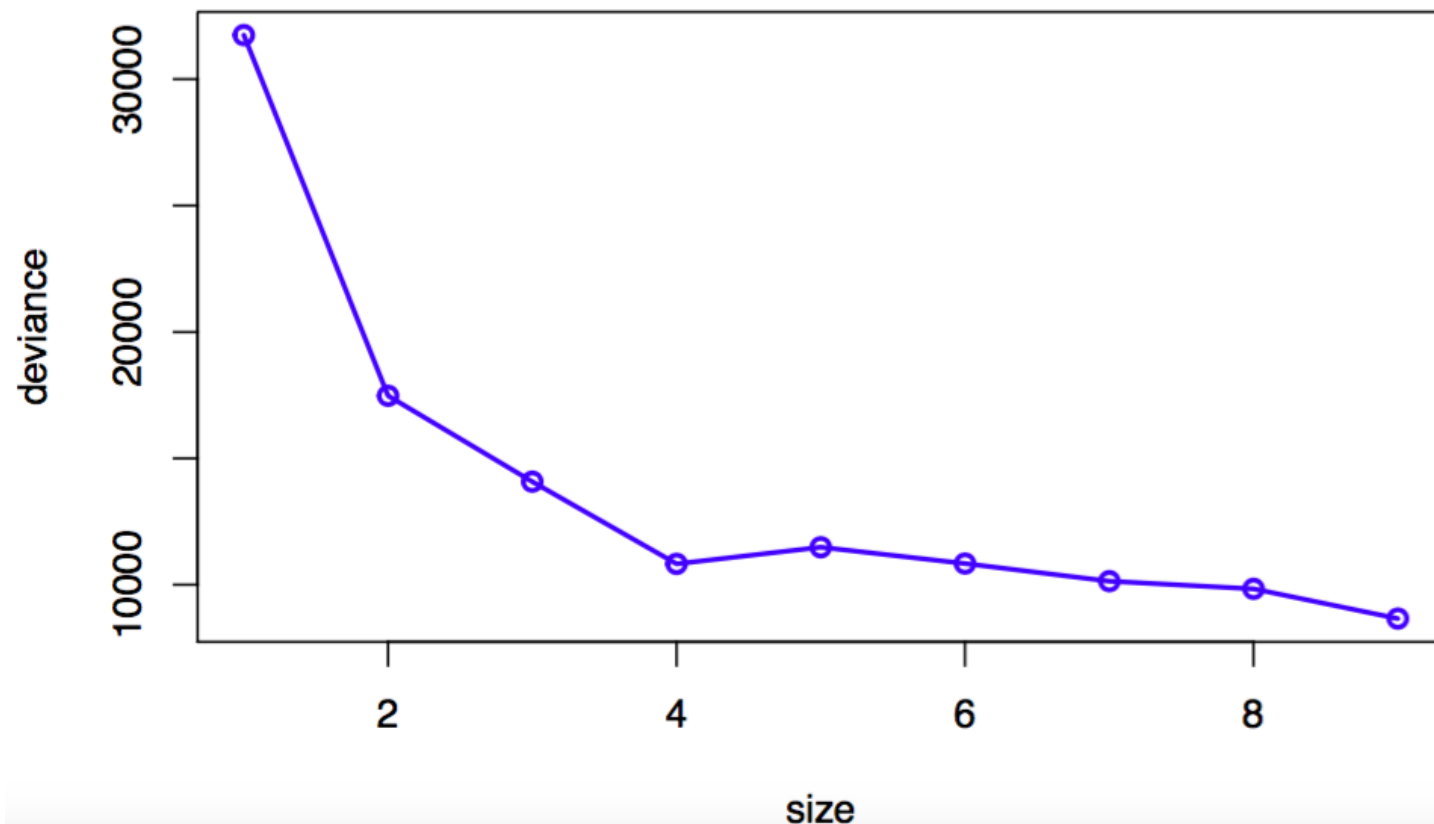


tree – cross-validation

```
cv.boston <- cv.tree(tree.boston)
```

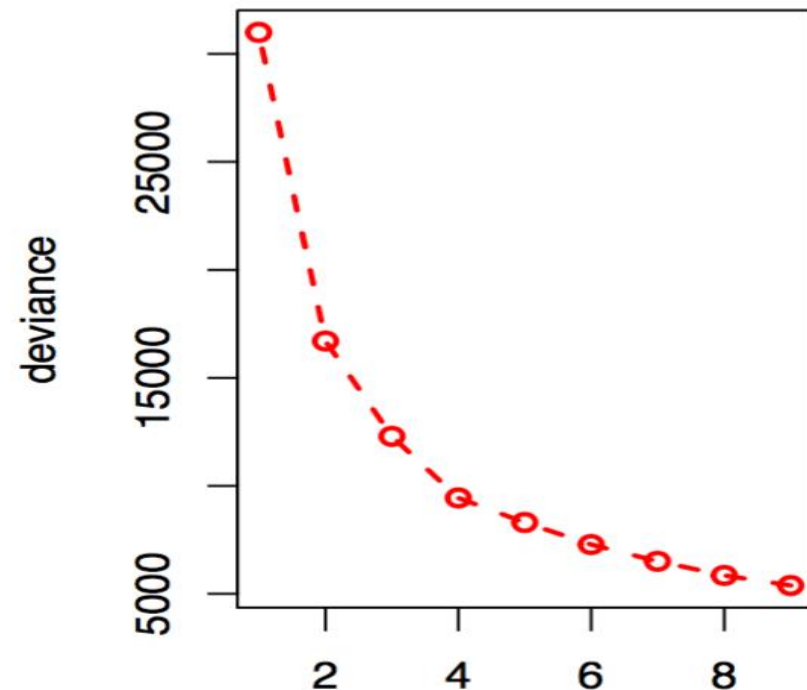
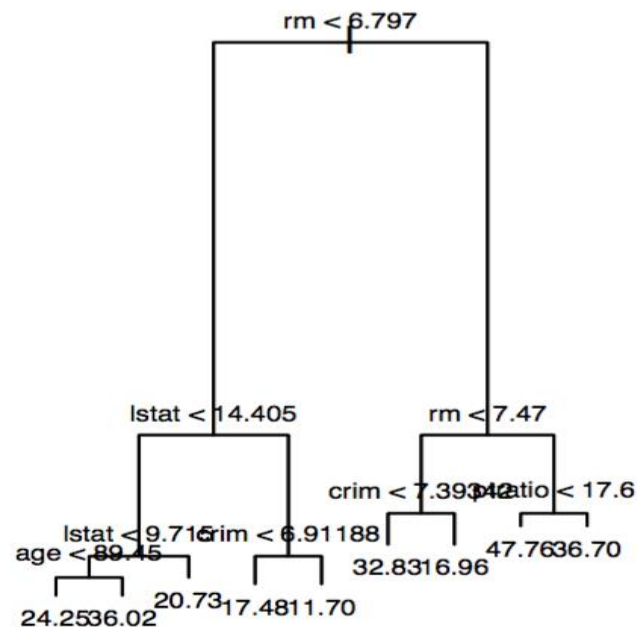
```
plot(cv.boston$size, cv.boston$dev, type='o', col=4, lwd=2, main="size vs. deviance")
```

size vs. deviance for cv



tree – pruning

```
prune.tree <- prune.tree(tree.boston, best=9)
```

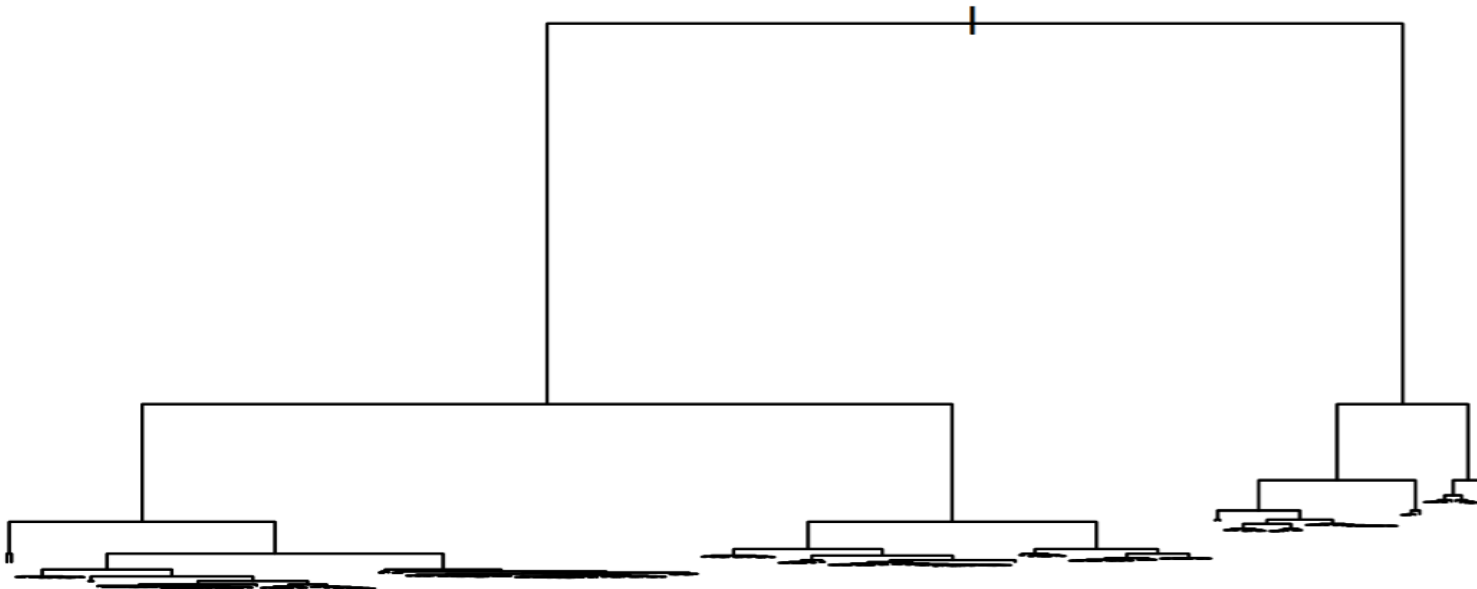


```
predict.tree <- predict(prune.tree, newdata=boston.test)
RMSE(predict.tree, y.true)
```

```
## [1] 4.262707
```

Pre-Pruning - tree

- mincut (min # obs. per node)
- minsize (max # obs. per node)
- mindev (within-node deviance must be at least this times that of the root for the node to be split)
- overfitting tree:

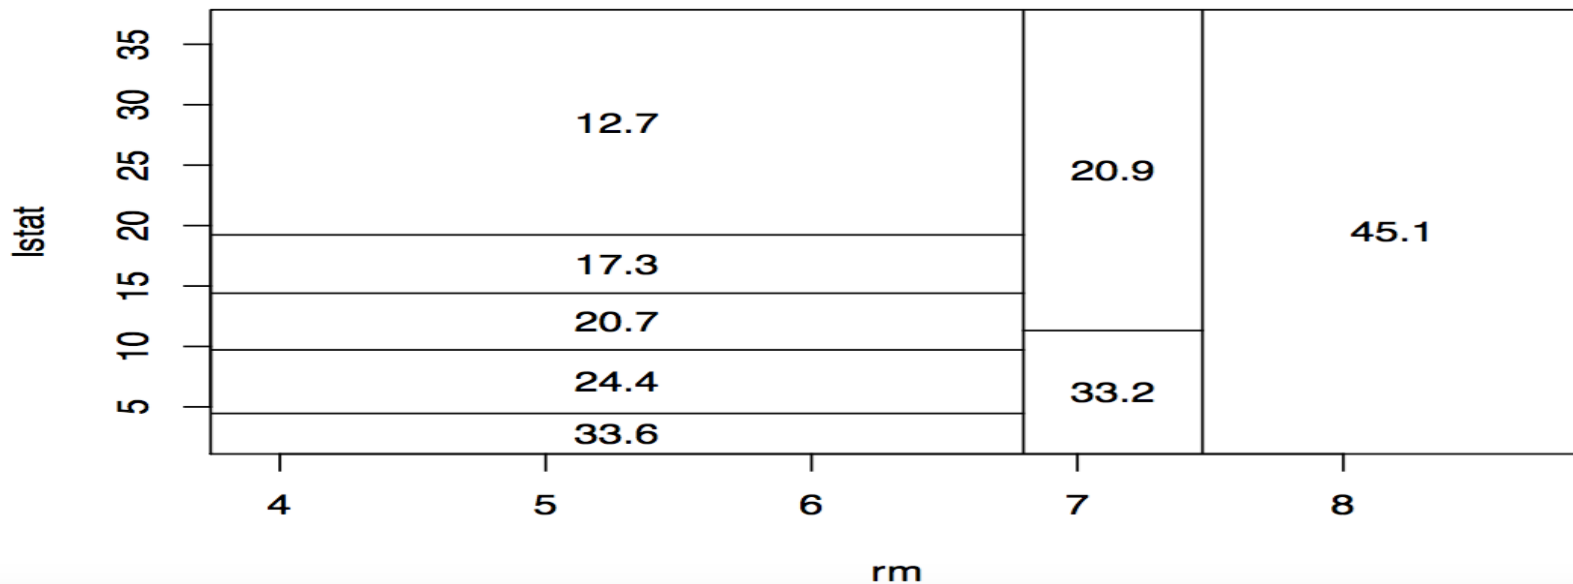


Pre-Pruning - rpart

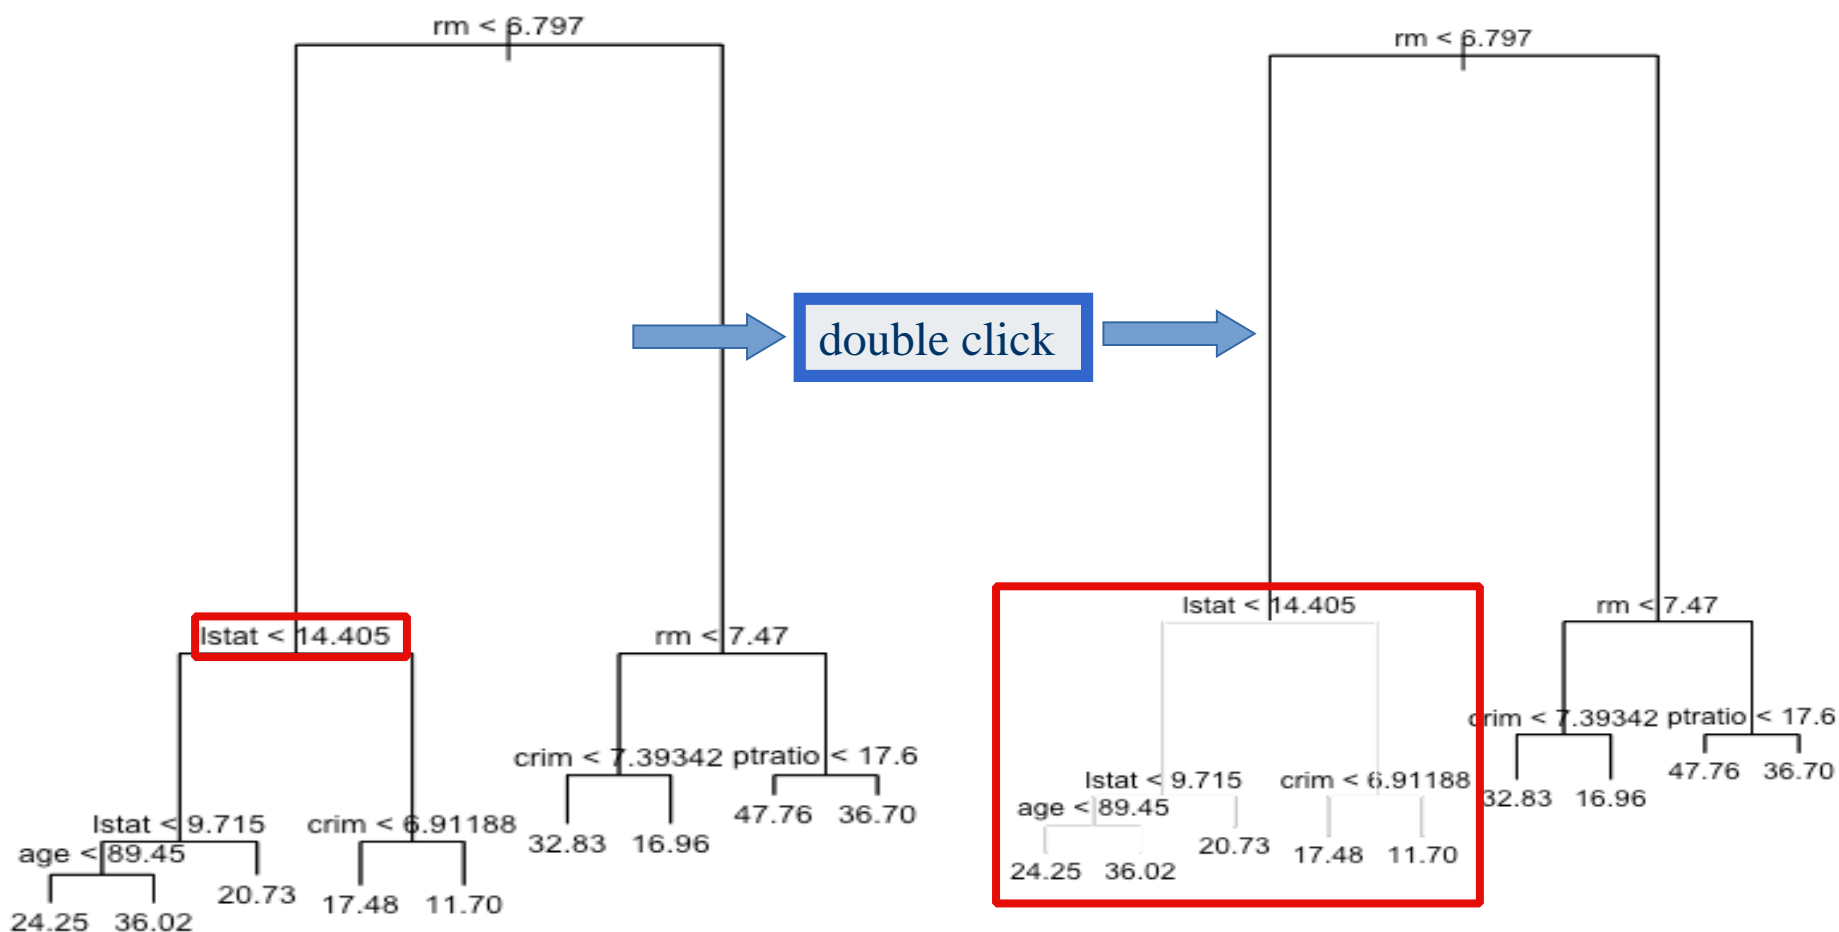
- minsplit (min # obs. per node)
- minbucket (max # obs. per node)
- maxdepth (of the tree)
- no check on deviance, but yes on depth

“partition.tree”

```
keep <- c("lstat", "rm", "medv")  
data2 <- data[, names(data) %in% keep]  
nrow(data2)  
  
## [1] 506  
names(data2)  
  
## [1] "rm"      "lstat"   "medv"  
new.tree <- tree(medv~., data2, subset=train)  
partition.tree(new.tree)
```



“snip.tree”



Handling Missing Values

rpart:

- surrogate variables

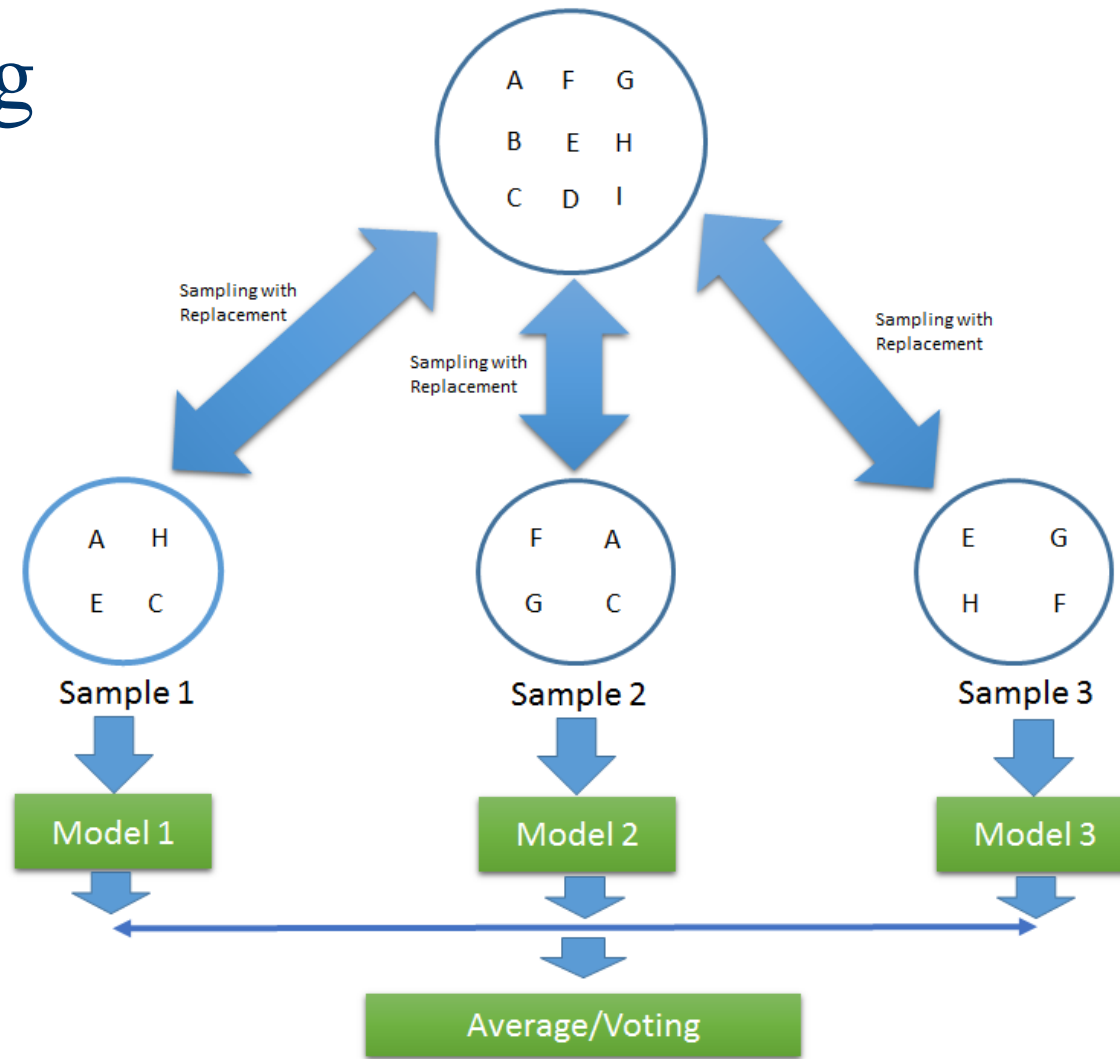
tree:

- drop as far as possible in the tree

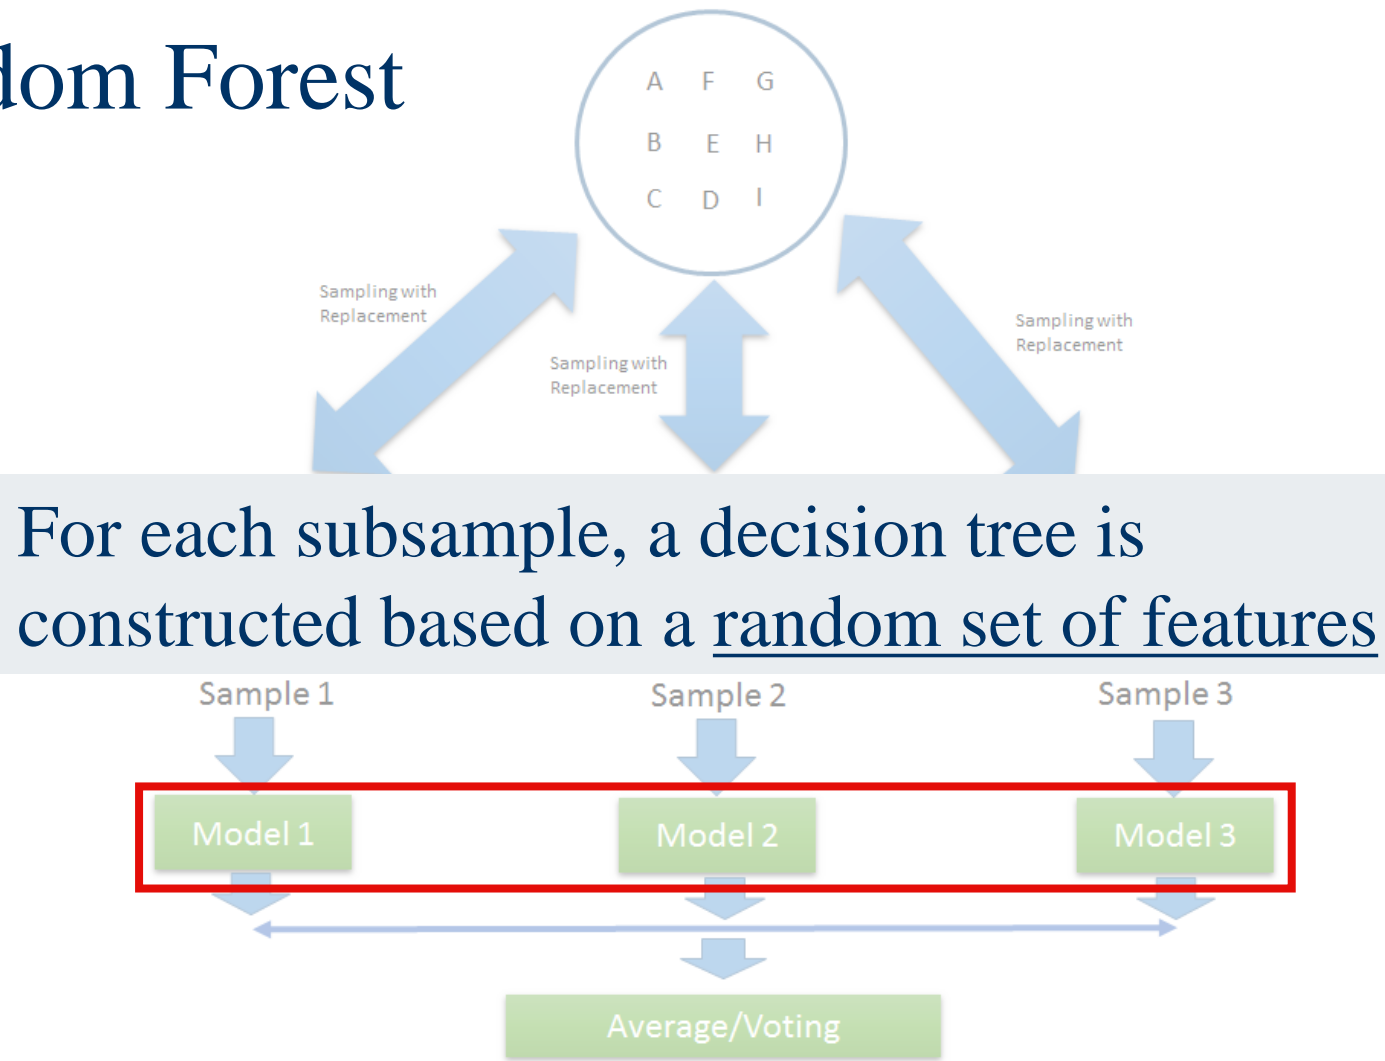
In any case, see “na.action” argument

Bagging and Random Forest in R

Bagging



Random Forest



Random Forest

Decorrelates the trees in the forest

Bagging x Random Forest

The main difference is the features considered in each tree

randomForest R package

Bagging
Random Forest

Regression
Classification

Package 'randomForest'

March 25, 2018

Title Breiman and Cutler's Random Forests for Classification and Regression

Version 4.6-14

Date 2018-03-22

Depends R (>= 3.2.2), stats

Suggests RColorBrewer, MASS

Author Fortran original by Leo Breiman and Adele Cutler, R port by Andy Liaw and Matthew Wiener.

Description Classification and regression based on a forest of trees using random inputs, based on Breiman (2001) <DOI:10.1023/A:1010933404324>.

Maintainer Andy Liaw <andy_liaw@merck.com>

License GPL (>= 2)

URL <https://www.stat.berkeley.edu/~breiman/RandomForests/>

NeedsCompilation yes

Repository CRAN

Date/Publication 2018-03-25 15:00:24 UTC

R topics documented:

classCenter	2
combine	3
getTree	4
grow	5
importance	6
importance85	7
margin	8
MDSplot	9
na.roughfix	10
outlier	11
partialPlot	12
plot.randomForest	14
predict.randomForest	15

randomForest package - Regression

Boston data set

```
library(randomForest)
library(MASS)
library(dplyr)
set.seed(2019)
```

```
# data
data <- Boston
glimpse(data)
```

```
# Train set
train = sample(1: nrow(data), nrow(data)*0.75)
```

```
# Test set
boston.test = data[-train, "medv"]
```

```
## Observations: 506
## Variables: 14
```

randomForest package - Regression

Random Forest

```
rand.boston =randomForest(medv~.,data ,subset =train,
                           ntree=1000, importance =TRUE)
# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE,          s'
##                Type of random forest: regression
##                Number of trees: 1000
## No. of variables tried at each split: 4
##
##                Mean of squared residuals: 12.18053
##                % Var explained: 85.1
```

randomForest package - Regression

Random Forest

```
rand.boston = randomForest(medv~.,data ,subset =train,
                           ntree=1000, importance =TRUE)

# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE, s'
##              Type of random forest: regression
##              Number of trees: 1000
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 12.18053
##              % Var explained: 85.1
```

randomForest package - Regression

Random Forest

```
rand.boston =randomForest(medv~.,data ,subset =train,
                           ntree=1000, importance =TRUE)
# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE,          s'
##              Type of random forest: regression
##              Number of trees: 1000
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 12.18053
##              % Var explained: 85.1
```

randomForest package - Regression

Random Forest

```
rand.boston =randomForest(medv~.,data ,subset =train,
                           ntree=1000 importance =TRUE)
# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE, s'
##               Type of random forest: regression
##               Number of trees: 1000
## No. of variables tried at each split: 4
##
##               Mean of squared residuals: 12.18053
##               % Var explained: 85.1
```


randomForest package - Regression

Random Forest

```
rand.boston =randomForest(medv~.,data ,subset =train,
                           ntree=1000, importance =TRUE)
# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE, s'
##               Type of random forest: regression
##               Number of trees: 1000
##               No. of variables tried at each split: 4
##
##               Mean of squared residuals: 12.18053
##               % Var explained: 85.1
```

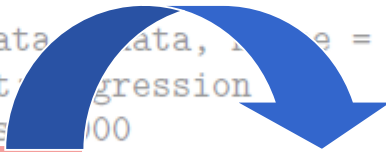
randomForest package - Regression

```

rand.boston =randomForest(medv~.,data ,subset =train,
                           ntree=1000, importance =TRUE)
# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE,
##               Type of random forest: regression
##               Number of trees: 1000
##               No. of variables tried at each split: 4
##
##               Mean of squared residuals
##               % Var explained

```



Here we have:

$13/3 \approx 4$ variables, since it is a Random Forest

randomForest package - Regression



```
rand.boston = randomForest(formula = medv ~ ., data = train,
                             ntree=1000, importance = TRUE)
```

```
# mtry: we chose the default (p/3)
# Number of variables randomly sampled at each split
# ntree: Number of trees to grow.
rand.boston
```

```
##
## Call:
## randomForest(formula = medv ~ ., data = train,
##               type = "rf", ntree = 1000,
##               importance = TRUE,
##               mtry = 13,
##               nodesize = 1,
##               ...)
```

Mean of squared residuals: 12.18053

% Var explained: 85.1

If we wanted to use the "Bagging" method we needed to add:

$mtry = (\text{total number of explanatory variables})$

In this case:

$mtry = 13$

randomForest package - Regression

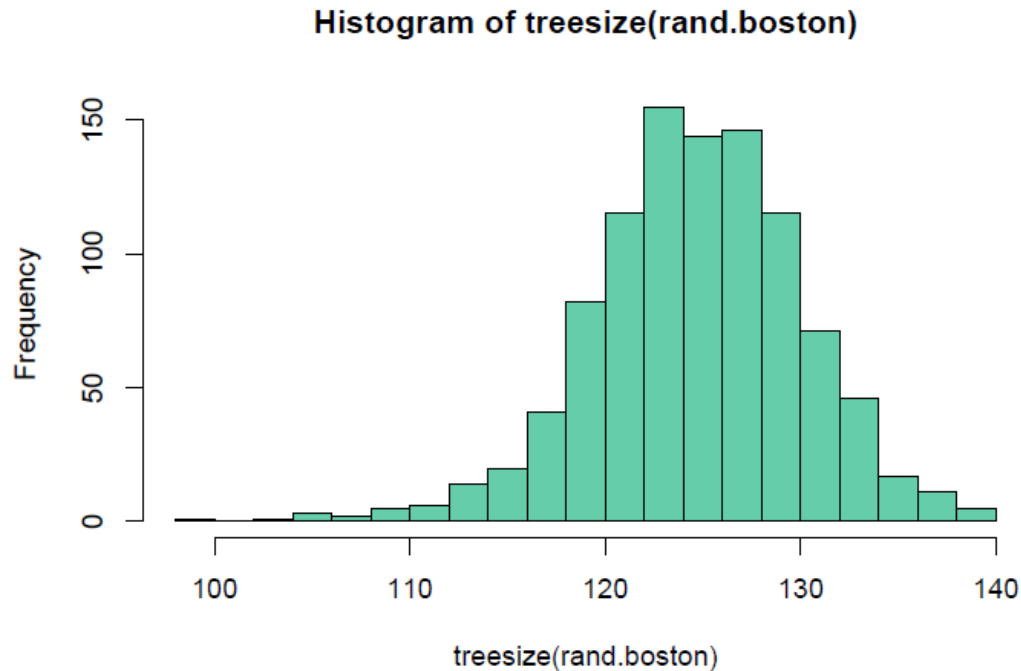
Random Forest

```
rand.boston =randomForest(medv~.,data ,subset =train,
                           ntree=1000, importance =TRUE)
# mtry: we chose the default (p/3)
#Number of variables randomly sampled as candidates at each split.
# ntree: Number of trees to grow.
rand.boston

##
## Call:
##  randomForest(formula = medv ~ ., data = data, ntree = 1000, importance = TRUE, s'
##              Type of random forest: regression
##              Number of trees: 1000
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 12.18053
##              % Var explained: 85.1
```

randomForest package - Regression

```
# Histogram of trees size  
hist(treesize(rand.boston), breaks=20, col="aquamarine3")
```

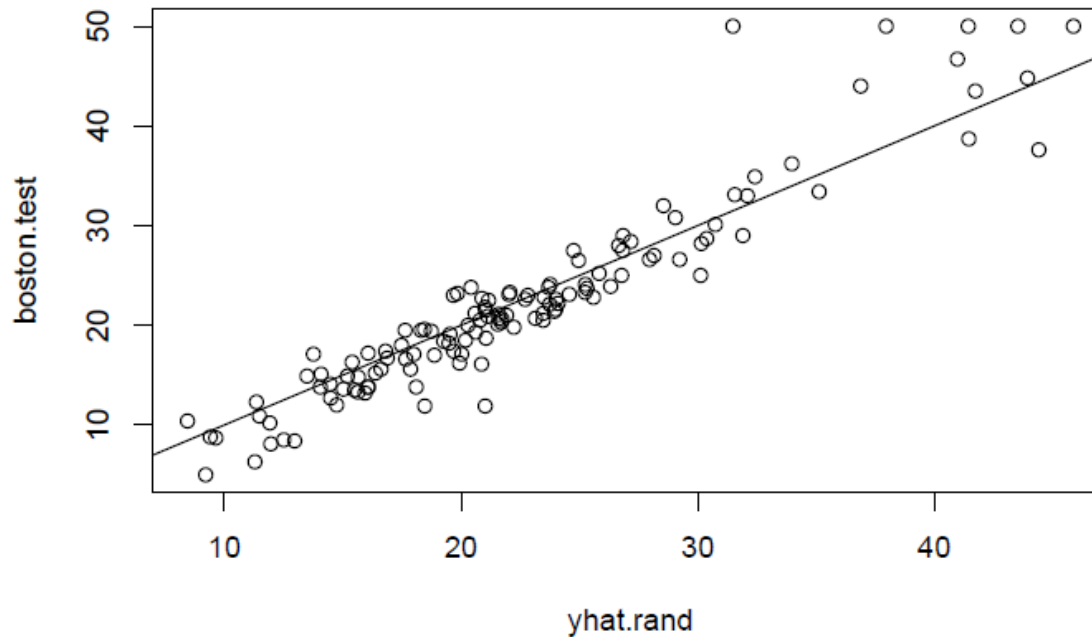


tree size = number of nodes

randomForest package - Regression

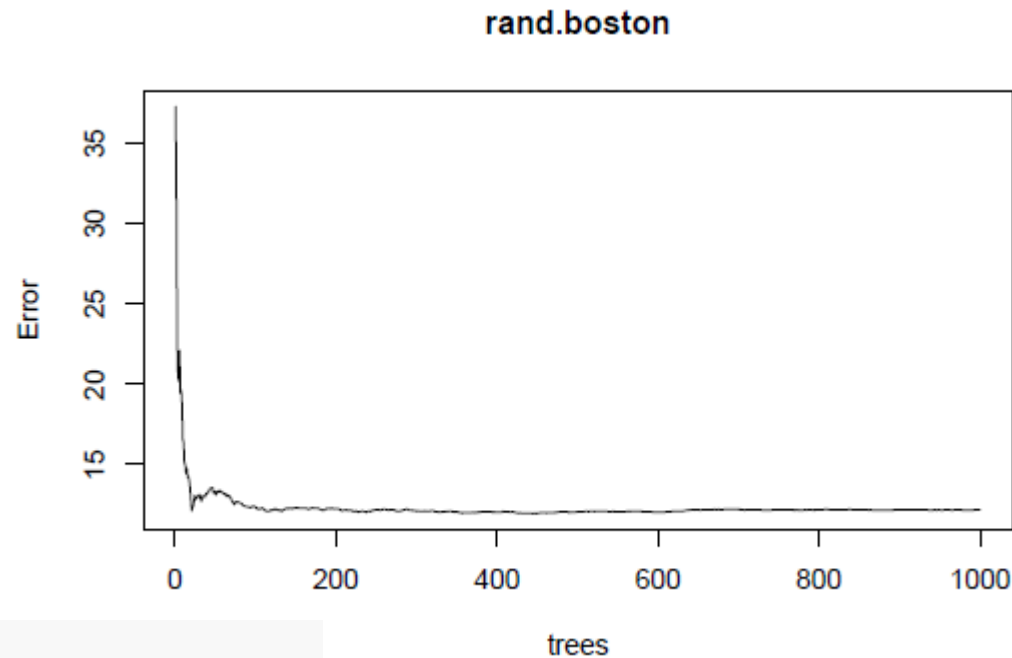
```
# Predictions
```

```
yhat.rand = predict (rand.boston , newdata = data[-train ,])  
plot(yhat.rand , boston.test)  
abline (0,1)
```



randomForest package - Regression

```
# Plot the error rates or MSE of a randomForest object x number of trees  
plot(rand.boston, type="l")
```

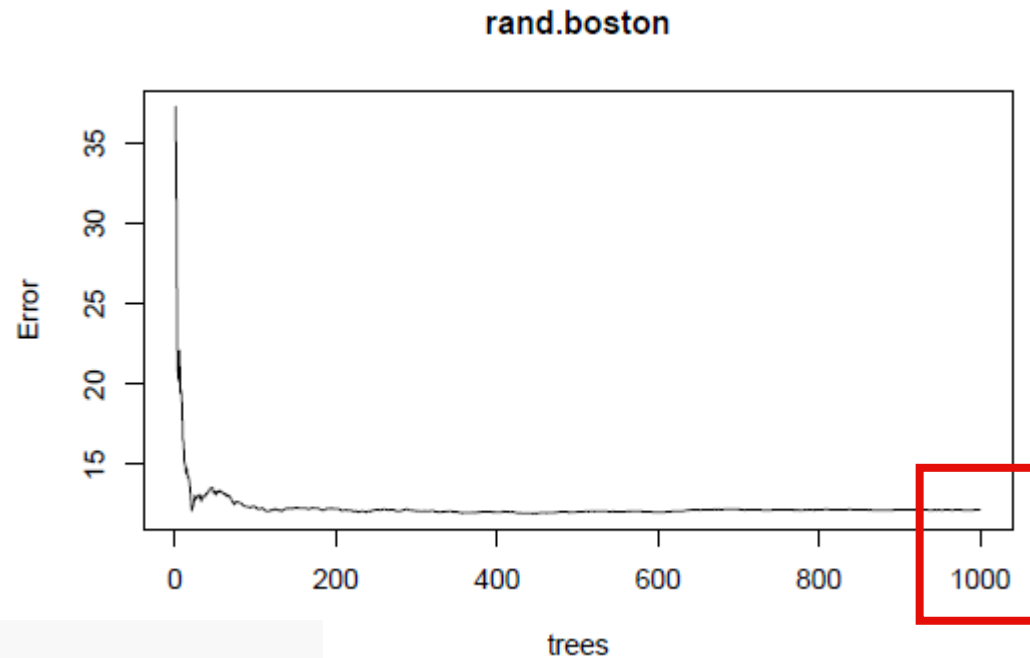


```
# MSE  
mean(( yhat.rand -boston.test)^2)
```

```
## [1] 10.84324
```

randomForest package - Regression

```
# Plot the error rates or MSE of a randomForest object x number of trees  
plot(rand.boston, type="l")
```



```
# MSE  
mean(( yhat.rand -boston.test)^2)
```

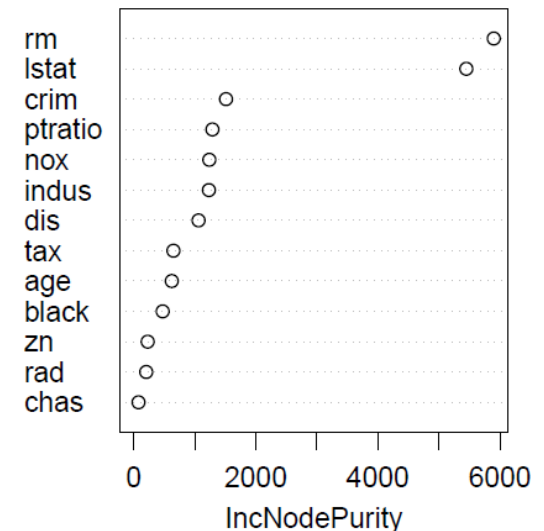
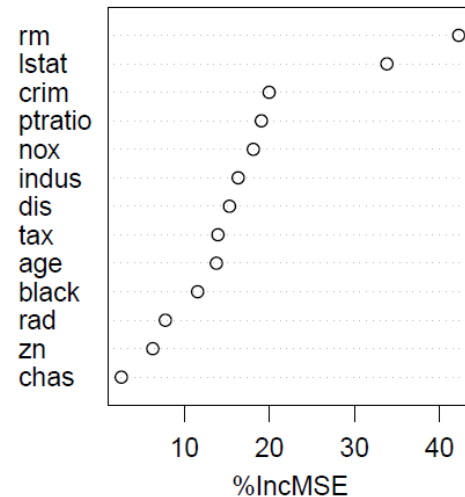
```
## [1] 10.84324
```


randomForest package - Regression

```
# Importance of variables
importance (rand.boston)
```

##		%IncMSE	IncNodePurity
##	crim	21.378868	1849.6242
##	zn	5.327068	327.3743
##	indus	15.023468	2064.4995
##	chas	4.422907	161.2110
##	nox	22.588932	1997.7708
##	rm	48.005356	8812.4332
##	age	18.105861	866.5470
##	dis	21.759105	1784.2645
##	rad	7.799535	253.0358
##	tax	17.556229	1116.2127
##	ptratio	22.875649	2033.8358
##	black	11.845736	577.0695
##	lstat	38.052427	8568.2762

```
varImpPlot (rand.boston)
```



randomForest package - Classification

Iris data set

```
library(randomForest)
library(caret)
library(dplyr)
set.seed(2019)
```

```
# data
data <- iris
glimpse(iris)
```

```
# Train set
train = sample(1: nrow(data), nrow(data)*0.75)
```

```
# Test set
iris.test <- data[-train, "Species"]
```

```
## Observations: 150
## Variables: 5
```

randomForest package - Classification

Random Forest

```
rand.iris = randomForest(Species~.,data, subset=train,
                          ntree=1000, proximity=TRUE)
# we chose the default number of variables (sqrt(p))
rand.iris

##
## Call:
## randomForest(formula = Species ~ ., data = data, ntree = 1000,      proximity = TRUE,
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
## OOB estimate of error rate: 6.25%
## Confusion matrix:
##      setosa versicolor virginica class.error
## setosa      40         0         0 0.00000000
## versicolor   0        31         4 0.11428571
## virginica    0         3        34 0.08108108
```

randomForest package - Classification

Random Forest

```
rand.iris = randomForest(Species~.,data, subset=train,
                        ntree=1000, proximity=TRUE)
# we chose the default number of variables (sqrt(p))
rand.iris

##
## Call:
## randomForest(formula = Species ~ ., data = data, ntree = 1000,      proximity = TRUE,
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
## OOB estimate of error rate: 6.25%
## Confusion matrix:
##      setosa versicolor virginica class.error
## setosa      40         0         0 0.00000000
## versicolor   0        31         4 0.11428571
## virginica    0         3        34 0.08108108
```

randomForest package - Classification

Random Forest

```

rand.iris = randomForest(Species~.,data, subset=train,
                          ntree=1000, proximity=TRUE)
# we chose the default number of variables (sqrt(p))
rand.iris

##
## Call:
##  randomForest(formula = Species ~ ., data = data, ntree = 1000,      proximity = TRUE,
##               Type of random forest: classification
##               Number of trees: 1000
## No. of variables tried at each split: 2
##
##           OOB estimate of  error rate: 6.25%
## Confusion matrix:
##           setosa versicolor virginica class.error
## setosa      40          0          0 0.00000000
## versicolor   0         31          4 0.11428571
## virginica    0          3         34 0.08108108

```

randomForest package - Classification


Random Forest

```

rand.iris = randomForest(Species~.,data, subset=train,
                          ntree=1000, proximity=TRUE)
# we chose the default number of variables (sqrt(p))
rand.iris

##
## Call:
## randomForest(formula = Species ~ ., data = data, ntree = 1000, proximity = TRUE,
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
## OOB estimate of error rate:
## Confusion matrix:
##      setosa versicolor virginica
## setosa      40         0         0 0.00000000
## versicolor   0        31         4 0.11428571
## virginica    0         3        34 0.08108108

```



Here we have:

$\sqrt{4} = 2$ variables, since it is a Random Forest

randomForest package - Classification

Random Forest

```
rand.iris = randomForest(Species~.,data, subset=train,
                          ntree=1000, proximity=TRUE)
# we chose the default number of variables (sqrt(p))
rand.iris

##
## Call:
## randomForest(formula = Species ~ ., data = data, ntree = 1000,      proximity = TRUE,
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
## OOB estimate of error rate: 6.25%
## Confusion matrix:
##           setosa versicolor virginica class.error
## setosa      40          0          0 0.00000000
## versicolor   0         31          4 0.11428571
## virginica    0          3         34 0.08108108
```

randomForest package - Classification

Random Forest

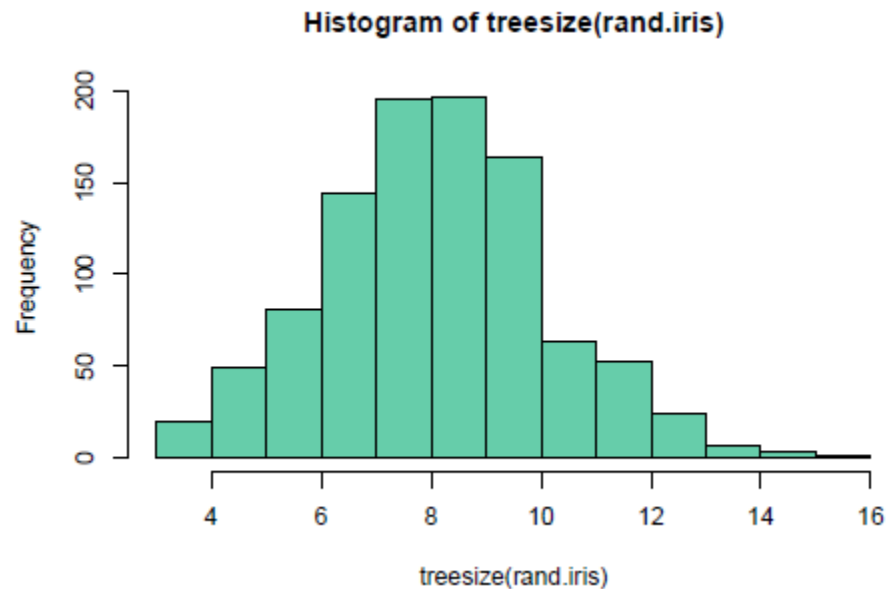
```
rand.iris = randomForest(Species~.,data, subset=train,
                          ntree=1000, proximity=TRUE)
# we chose the default number of variables (sqrt(p))
rand.iris
```

```
##
## Call:
## randomForest(formula = Species ~ ., data = data, ntree = 1000,      proximity = TRUE,
##              Type of random forest: classification
##              Number of trees: 1000
## No. of variables tried at each split: 2
##
##          OOB estimate of  error rate: 6.25%
## Confusion matrix:
```

	setosa	versicolor	virginica	class.error
setosa	40	0	0	0.00000000
versicolor	0	31	4	0.11428571
virginica	0	3	34	0.08108108

randomForest package - Classification

```
# Histogram of trees size  
hist(treesize(rand.iris), breaks=10, col="aquamarine3")
```

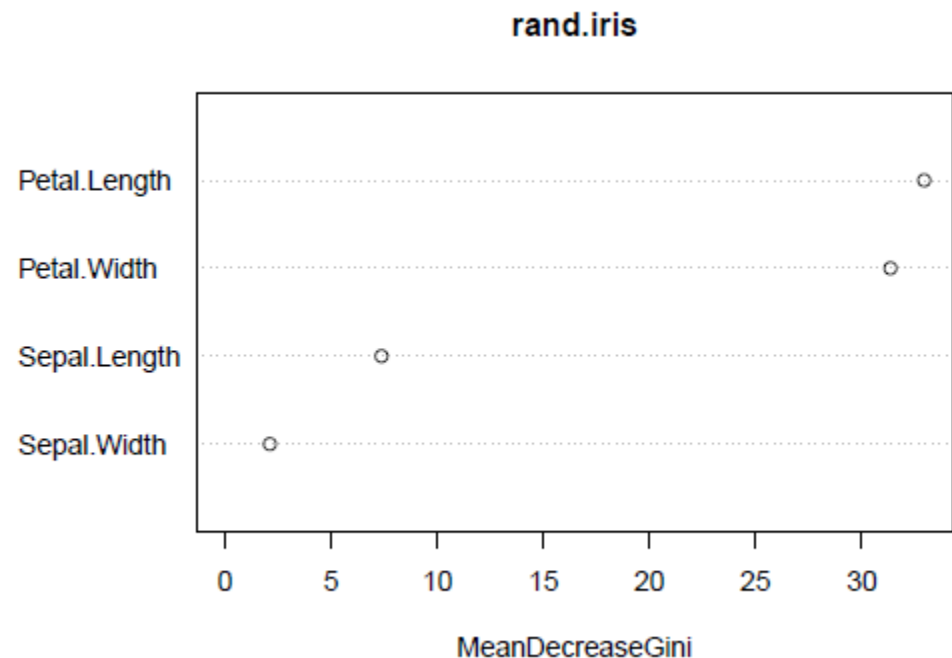


randomForest package - Classification

```
# Importance of variables  
importance(rand.iris)
```

```
##              MeanDecreaseGini  
## Sepal.Length      7.380465  
## Sepal.Width       2.121128  
## Petal.Length     32.945481  
## Petal.Width      31.354294
```

```
varImpPlot(rand.iris)
```



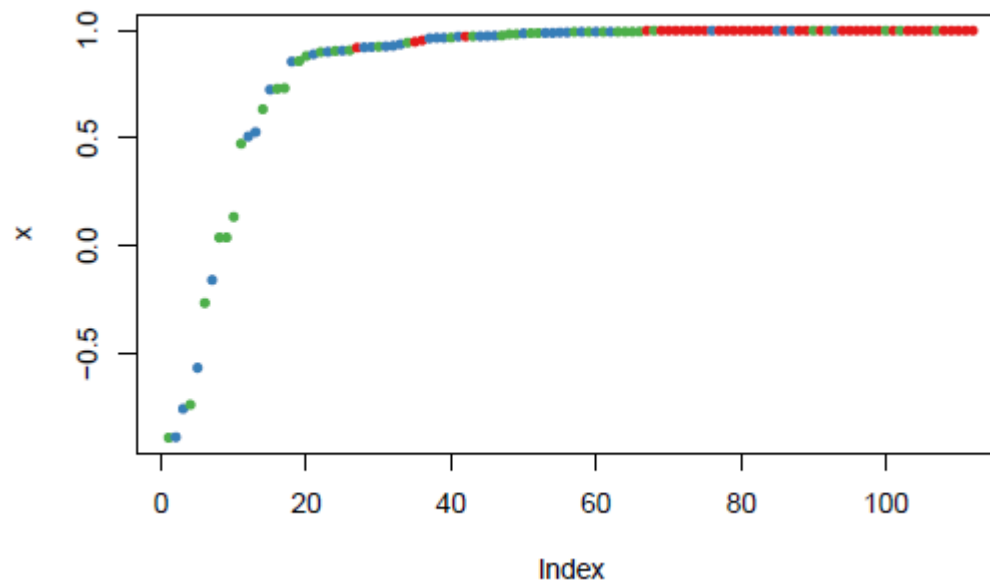
randomForest package - Classification

```
# Prediction  
irisPred <- predict(rand.iris , newdata = data[-train ,])  
table(irisPred, iris.test)
```

```
##           iris.test  
## irisPred   setosa versicolor virginica  
##   setosa      10         0         0  
## versicolor   0         15         0  
##   virginica   0         0        13
```

randomForest package - Classification

```
# Margin  
# Compute or plot the margin of predictions from a randomForest classifier  
margin.iris <- randomForest::margin(rand.iris,iris.test)  
plot(margin.iris)
```



The margin of a data point is defined as the proportion of votes for the correct class minus maximum proportion of votes for the other classes. Thus under majority votes, positive margin means correct classification, and vice versa.

randomForest package - Classification

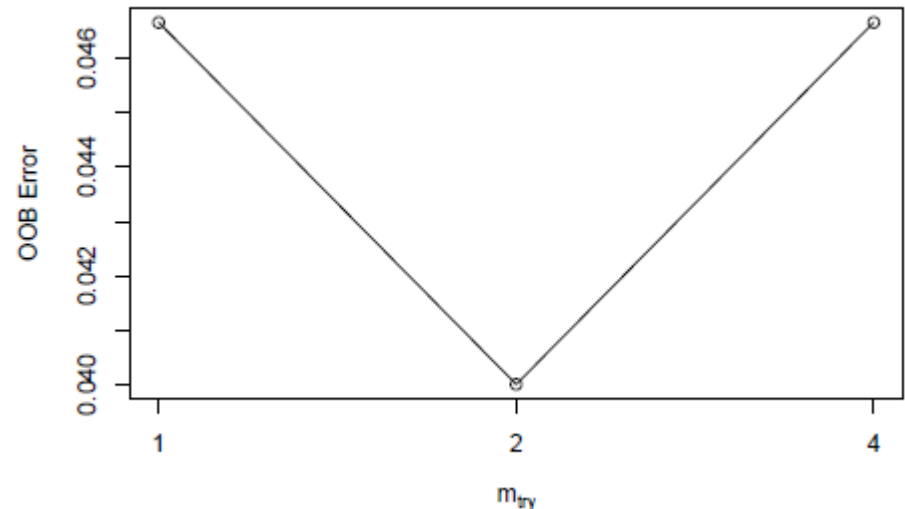
```
# Tuning
# Tune randomForest for the optimal mtry parameter
tune.rf <- tuneRF(iris[,-5], iris[,5], stepFactor=0.5)
```

```
## mtry = 2  OOB error = 4%
## Searching left ...
## mtry = 4    OOB error = 4.67%
## -0.1666667 0.05
## Searching right ...
## mtry = 1    OOB error = 4.67%
## -0.1666667 0.05
```

```
tune.rf
```

```
##      mtry  OOBError
## 1.00B    1 0.04666667
## 2.00B    2 0.04000000
## 4.00B    4 0.04666667
```

Feature Selection



Generalized Boosted Regression Models

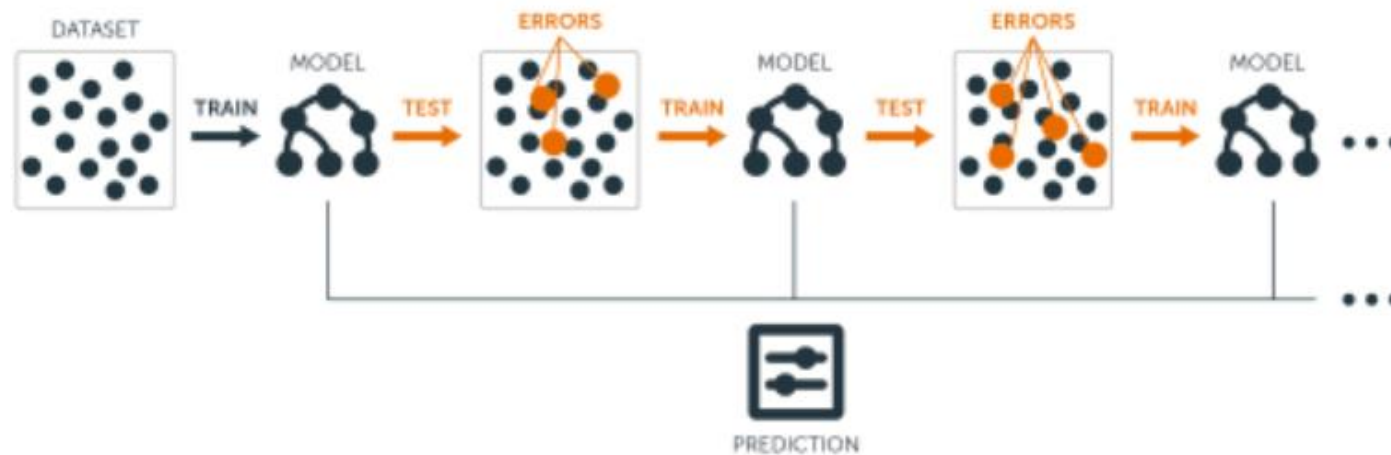
gbm package

Generalized Boosted Regression Models

gbm package

- Gradient Boosting trains many models in a gradual, additive and sequential manner.
- The major difference between AdaBoost and Gradient Boosting Algorithm is how the two algorithms identify the shortcomings of weak learners (eg. decision trees).
- While the AdaBoost model identifies the shortcomings by using high weight data points, gradient boosting performs the same by using gradients in the loss function

- The main idea of boosting is to add new models to the ensemble *sequentially*. At each particular iteration, a new weak, base-learner model is trained with respect to the error of the whole ensemble learnt so far.



Boosting has three tuning parameters

- **Number of trees:** The total number of trees to fit. GBMs often require many trees; however, unlike random forests GBMs can overfit so the goal is to find the optimal number of trees that minimize the loss function of interest with cross validation.
- **Depth of trees:** The number d of splits in each tree, which controls the complexity of the boosted ensemble.
- **Learning rate:** The shrinkage parameter. It controls the rate in which boosting learns.

Package ‘gbm’

January 14, 2019

Version 2.1.5

Title Generalized Boosted Regression Models

Depends R (>= 2.9.0)

Imports gridExtra, lattice, parallel, survival

Suggests knitr, pdp, RUnit, splines, viridis

Description An implementation of extensions to Freund and Schapire's AdaBoost algorithm and Friedman's gradient boosting machine. Includes regression methods for least squares, absolute loss, t-distribution loss, quantile regression, logistic, multinomial logistic, Poisson, Cox proportional hazards partial likelihood, AdaBoost exponential loss, Huberized hinge loss, and Learning to Rank measures (LambdaMart). Originally developed by Greg Ridgeway.

License GPL (>= 2) | file LICENSE

URL <https://github.com/gbm-developers/gbm>

BugReports <https://github.com/gbm-developers/gbm/issues>

RoxygenNote 6.1.1

VignetteBuilder knitr

NeedsCompilation yes

Author Brandon Greenwell [aut, cre] (<<https://orcid.org/0000-0002-8120-0084>>),
Bradley Boehmke [aut] (<<https://orcid.org/0000-0002-3611-8516>>),
Jay Cunningham [aut],
GBM Developers [aut] (<https://github.com/gbm-developers>)

Maintainer Brandon Greenwell <greenwell1.brandon@gmail.com>

Repository CRAN

Date/Publication 2019-01-14 15:00:03 UTC

R topics documented:

gbm-package	2
basehaz.gbm	3
calibrate.plot	4

The gbm R package is an implementation of extensions to Freund and Schapire's AdaBoost algorithm and Friedman's gradient boosting machine. This is the original R implementation of GBM.

URL <https://github.com/gbm-developers/gbm>

gbm

Generalized Boosted Regression Modeling (GBM)

Description

Fits generalized boosted regression models. For technical details, see the vignette: `utils::browseVignettes("gbm")`.

Usage

```
gbm(formula = formula(data), distribution = "bernoulli",  
    data = list(), weights, var.monotone = NULL, n.trees = 100,  
    interaction.depth = 1, n.minobsinnode = 10, shrinkage = 0.1,  
    bag.fraction = 0.5, train.fraction = 1, cv.folds = 0,  
    keep.data = TRUE, verbose = FALSE, class.stratify.cv = NULL,  
    n.cores = NULL)
```

Arguments:

- `formula` - A symbolic description of the model to be fit
- `distribution` - name of the distribution
- `n.trees` - Integer specifying the total number of trees to fit
- `interaction.depth` - the maximum depth of each tree
- `shrinkage` - the learning rate or step-size reduction

```
train <- sample(1: nrow(Boston), nrow(Boston) * 0.75)
# Test and Train set
Boston.train <- Boston[train,]
Boston.test <- Boston[-train,]
```

```
# Model for the Regression problem
# Default settings in gbm
boston.boost1<- gbm(medv~., data = Boston.train, cv.folds = 5)

## Distribution not specified, assuming gaussian ...
print(boston.boost1)

## gbm(formula = medv ~ ., data = Boston.train, cv.folds = 5)
## A gradient boosted model with gaussian loss function.
## 100 iterations were performed.
## The best cross-validation iteration was 99.
## There were 13 predictors of which 10 had non-zero influence.
```

summary.gbm

Summary of a gbm object

Description

Computes the relative influence of each variable in the gbm object.

Usage

```
## S3 method for class 'gbm'  
summary(object, cBars = length(object$var.names),  
  n.trees = object$n.trees, plotit = TRUE, order = TRUE,  
  method = relative.influence, normalize = TRUE, ...)
```

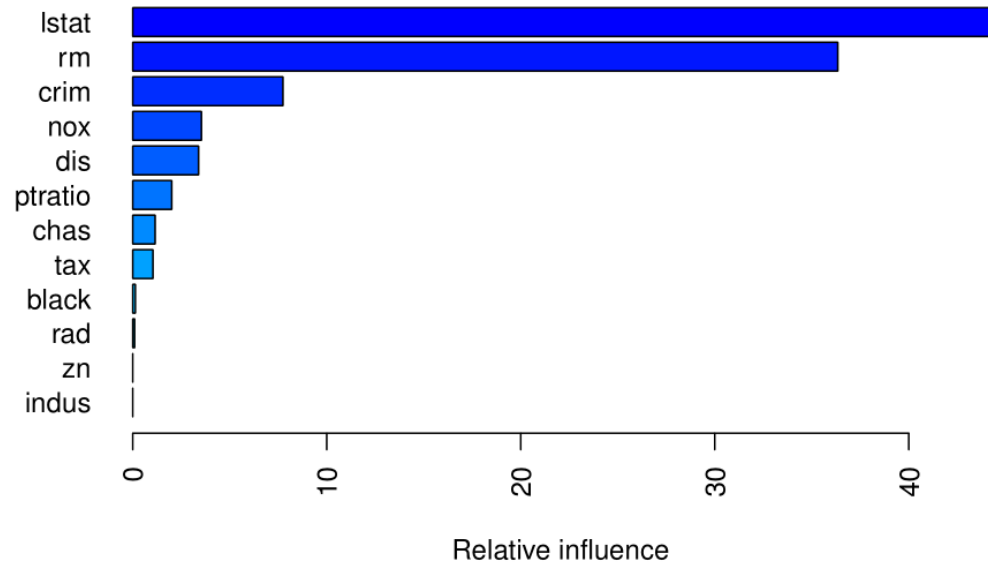
Arguments:

- object - a gbm object created from an initial call to gbm.
- cBars - number of bars to plot
- method - function used to compute the relative influence.

The permutation.test.gbm: Randomly permutes each predictor variable at a time and computes the associated reduction in predictive performance.

Variable importance

```
summary(boston.boost1, cBars = 12,
        method = relative.influence, # also can use permutation.test.gbm
        las = 2)
```



##	var	rel.inf
##	lstat	lstat 44.53540713
##	rm	rm 36.34203002
##	crim	crim 7.74309272
##	nox	nox 3.54144074
##	dis	dis 3.39437715
##	ptratio	ptratio 2.01057975
##	chas	chas 1.15397212
##	tax	tax 1.04375450
##	black	black 0.13674443
##	rad	rad 0.09860143
##	zn	zn 0.00000000
##	indus	indus 0.00000000
##	age	age 0.00000000

gbm.perf

GBM performance

Description

Estimates the optimal number of boosting iterations for a gbm object and optionally plots various performance measures

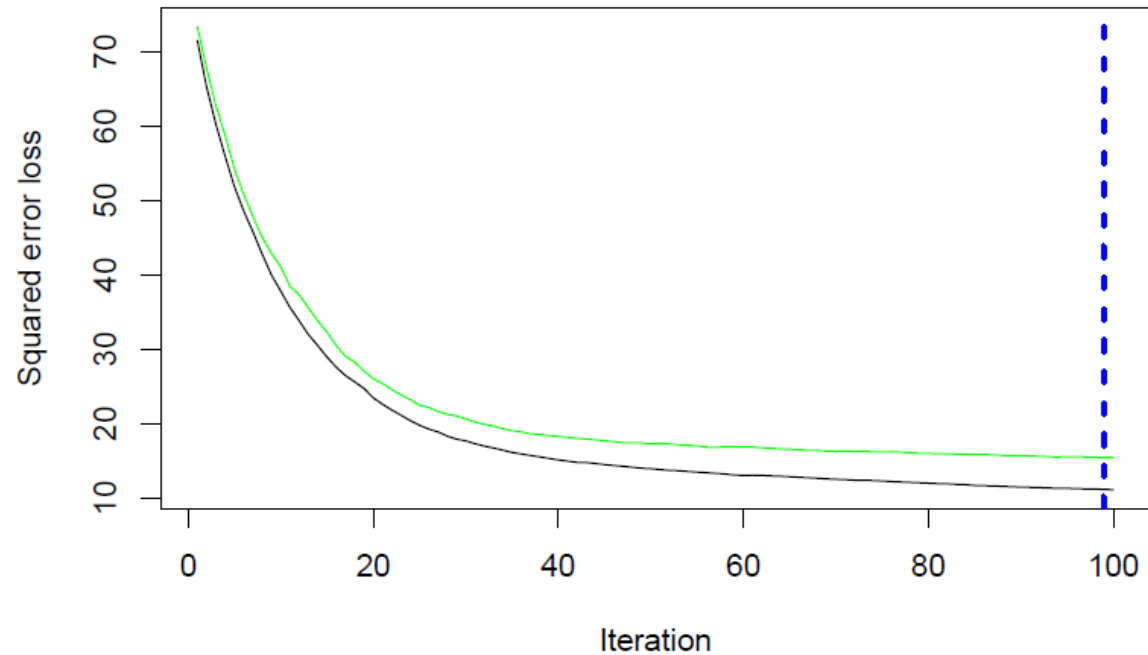
Usage

```
gbm.perf(object, plot.it = TRUE, oobag.curve = FALSE, overlay = TRUE,  
method)
```

Arguments:

- `method` - used to estimate the optimal number of boosting iterations.
`method = "OOB"` computes the out-of-bag estimate
`method = "cv"` extracts the optimal number of iterations using cross-validation if gbm was called with `cv.folds > 1`.

```
# plot loss function as a result of n trees added to the ensemble  
gbm.perf(boston.boost1, method = "cv")
```



```
## [1] 99
```


Tuning parameters

```
# Tuning
# Create hyperparameter grid
hyper_grid <- expand.grid(
  shrinkage = c(0.01, 0.02, 0.05, 0.08, 0.1),
  interaction.depth = c(3, 5, 8),
  optimal_trees = 0,           # a place to dump results
  min_RMSE = 0                # a place to dump results
)

# total number of combinations
nrow(hyper_grid)

## [1] 15
```

```
for(i in 1:nrow(hyper_grid)) {

  # Train GBM model
  gbm.tune <- gbm(formula = medv ~ ., distribution = "gaussian", data = Boston.train,
    n.trees = 5000, interaction.depth = hyper_grid$interaction.depth[i],
    shrinkage = hyper_grid$shrinkage[i], cv.folds = 5
  )

  # add min training error and trees to grid
  hyper_grid$optimal_trees[i] <- which.min(gbm.tune$cv.error)
  hyper_grid$min_RMSE[i] <- sqrt(min(gbm.tune$cv.error))
}
```

Tuning parameters

```
hyper_grid %>%  
dplyr::arrange(min_RMSE) %>%  
head(10)
```

##	shrinkage	interaction.depth	optimal_trees	min_RMSE
## 1	0.01	5	2323	3.444757
## 2	0.08	8	481	3.446742
## 3	0.01	3	3411	3.455078
## 4	0.08	5	251	3.456847
## 5	0.05	8	457	3.458843
## 6	0.10	3	226	3.512021
## 7	0.05	3	602	3.519723
## 8	0.02	8	904	3.536458
## 9	0.02	5	1642	3.548916
## 10	0.02	3	1058	3.549183

```
boston.opt <- gbm(formula = medv ~ .,distribution = "gaussian",data = Boston.train,  
  n.trees = 2323, interaction.depth = 5,shrinkage = 0.01,cv.folds = 5)
```

predict.gbm

Predict method for GBM Model Fits

Description

Predicted values based on a generalized boosted model object

Usage

```
## S3 method for class 'gbm'  
predict(object, newdata, n.trees, type = "link",  
        single.tree = FALSE, ...)
```

The predict.gbm produces predicted values for each observation in newdata using the first n.trees iterations of the boosting sequence.

```
#Predicting values for test data
```

```
boston.opt.test.pred.tree = predict(boston.opt, Boston.test, n.trees = boston.opt$n.trees)
MSE.test.boosting.opt <- mean((boston.opt.test.pred.tree - Boston.test$medv)^2)
MSE.test.boosting.opt
```

```
## [1] 9.335418
```

```
library(MLmetrics) # for MSE
```

```
MSE(boston.opt.test.pred.tree, Boston.test$medv)
```

```
## [1] 9.335418
```

```
#Predicting values for train data
```

```
boston.opt.train.pred.tree = predict(boston.opt, n.trees = boston.opt$n.trees)  
MSE.train.boosting.opt <- mean((boston.opt.train.pred.tree - Boston.train$medv)^2)  
MSE.train.boosting.opt
```

```
## [1] 1.641844
```

```
MSE(boston.opt.train.pred.tree, Boston.train$medv)
```

```
## [1] 1.641844
```

Methods	MSE – Test Error
Tree	18.17
rpart	16.97
Random Forest	10.84
Boosting	9.34

Thank you very much for your attention

Grazie mille per la vostra attenzione

Muito obrigado pela atenção